

视听信息系统导论第二次编程作业报告

目录

1	实验任务	2
1.1	Task1	2
1.2	Task2	2
1.3	Task3	3
1.4	Task4	4
1.5	Task5	5
1.6	Task6	6
1.7	Task7	8
1.8	Task8	9

1 实验任务

1.1 Task1

完成 model.py 文件中 `__init__` 函数。(完成代码即可，不用在报告中写文字说明)

代码如下：

```
1  def __init__(self, in_dim=1280, hidden_dim=256, num_classes=20, \
2      roi_output_w=2, roi_output_h=2, drop_ratio=0.3):
3      super().__init__()
4
5      assert(num_classes != 0)
6      self.num_classes = num_classes
7      self.roi_output_w, self.roi_output_h = roi_output_w, roi_output_h
8      self.feats_extractor = FeatureExtractor()
9      #####
10     # TODO: Declare the cls & bbox heads (in Fast R-CNN).           #
11     # The cls & bbox heads share a sequential module with a Linear layer, #
12     # followed by a Dropout (p=drop_ratio), a ReLU nonlinearity and another #
13     # Linear layer.                                                    #
14     # The cls head is a Linear layer that predicts num_classes + 1 (background). #
15     # The det head is a Linear layer that predicts offsets(dim=4).      #
16     # HINT: The dimension of the two Linear layers are in_dim -> hidden_dim and #
17     # hidden_dim -> hidden_dim.                                         #
18     #####
19     # Replace "pass" statement with your code
20     self.shared_fc = nn.Sequential(
21         nn.Linear(in_dim, hidden_dim),
22         nn.Dropout(drop_ratio),
23         nn.ReLU(),
24         nn.Linear(hidden_dim, hidden_dim)
25     )
26     self.cls_head = nn.Linear(hidden_dim, num_classes + 1)
27     self.bbox_head = nn.Linear(hidden_dim, 4)
28     #####
29     #                               END OF YOUR CODE                               #
30     #####
```

1.2 Task2

完成 utils.py 文件中的 `compute_iou` 函数。(完成代码即可，不用在报告中写文字说明)

代码如下：

```
1  def compute_iou(anchors, bboxes):
2      """
3      Compute the intersection-over-union between anchors and gts.
4
5      Inputs:
6      - anchors: Anchor boxes, of shape (M, 4), where M is the number of proposals
7      - bboxes: GT boxes of shape (N, 4), where N is the number of GT boxes,
8          4 indicates (xlgt, ylgt, xrgt, yrgt)
```

```

9
10 Outputs:
11 - iou: IoU matrix of shape (M, N)
12 """
13 iou = None
14 #####
15 # TODO: Given anchors and gt bboxes, #
16 # compute the iou between each anchor and gt bbox. #
17 #####
18 x1 = torch.max(anchors[:, None, 0], bboxes[None, :, 0])
19 y1 = torch.max(anchors[:, None, 1], bboxes[None, :, 1])
20 x2 = torch.min(anchors[:, None, 2], bboxes[None, :, 2])
21 y2 = torch.min(anchors[:, None, 3], bboxes[None, :, 3])
22 inter = torch.clamp(x2 - x1, min=0) * torch.clamp(y2 - y1, min=0)
23 area_anchors = (anchors[:, 2] - anchors[:, 0]) * (anchors[:, 3] - anchors[:, 1])
24 area_bboxes = (bboxes[:, 2] - bboxes[:, 0]) * (bboxes[:, 3] - bboxes[:, 1])
25 union = area_anchors[:, None] + area_bboxes[None, :] - inter
26 iou = torch.zeros_like(union)
27 non_zero_union = union > 0
28 iou[non_zero_union] = inter[non_zero_union] / union[non_zero_union]
29 #####
30 # END OF YOUR CODE #
31 #####
32
33 return iou

```

1.3 Task3

阅读 `utils.py` 文件中的 `assign_label` 函数，并简要说明该函数如何判断正负样本框。

答：assign_label 函数用于在模型训练中将候选框（proposals）分配为正样本或负样本。如果某个候选框与任意一个 GT 框的 IoU 值是所有候选框中的最大值，或者当候选框与某个 GT 框的 IoU 大于正样本阈值（pos_thresh）时，则该候选框被标记为正样本。当候选框与所有 GT 框的 IoU 值都小于负样本阈值（neg_thresh）时，该候选框被标记为负样本，此外再随机采样负样本以平衡正负比例。

阅读 `utils.py` 文件中的 `compute_offsets` 函数，简要说明如何计算正样本框到真实框的偏移量。

答：compute_offsets 函数用于计算正样本框（anchors）相对于真实框（GT boxes）的偏移量，以便在模型训练中进行位置调整。

计算 GT 框和 anchor 框的宽高比例，通过 torch.log 函数取对数来得到宽高偏移量：

$$\Delta w = \log\left(\frac{w^{\text{gt}}}{w^{\text{anchor}}}\right), \quad \Delta h = \log\left(\frac{h^{\text{gt}}}{h^{\text{anchor}}}\right) \quad (1)$$

计算 GT 框与 anchor 框的中心坐标差异，并对 anchor 框的宽高进行归一化来得到中心坐标偏移量：

$$\Delta x = \frac{x_{\text{center}}^{\text{gt}} - x_{\text{center}}^{\text{anchor}}}{w^{\text{anchor}}}, \quad \Delta y = \frac{y_{\text{center}}^{\text{gt}} - y_{\text{center}}^{\text{anchor}}}{h^{\text{anchor}}} \quad (2)$$

随后将这四个偏移量拼接在一起，得到最终的偏移量。

$$\text{offsets} = (\Delta x, \Delta y, \Delta w, \Delta h) \quad (3)$$

1.4 Task4

完成 model.py 文件中的 forward 函数。(完成代码即可，不用在报告中写文字说明)

代码如下：

```

1  def forward(self, images, bboxes, bbox_batch_ids, proposals, proposal_batch_ids):
2      """
3      Training-time forward pass for our two-stage Faster R-CNN detector.
4
5      Inputs:
6      - images: Tensor of shape (B, 3, H, W) giving input images
7      - bboxes: Tensor of shape (N, 5) giving ground-truth bounding boxes
8      and category labels, from the dataloader, where N is the total number
9      of GT boxes in the batch
10     - bbox_batch_ids: Tensor of shape (N, ) giving the index (in the batch)
11     of the image that each GT box belongs to
12     - proposals: Tensor of shape (M, 4) giving the proposals for input images,
13     where M is the total number of proposals in the batch
14     - proposal_batch_ids: Tensor of shape (M, ) giving the index of the image
15     that each proposals belongs to
16
17     Outputs:
18     - total_loss: Torch scalar giving the overall training loss.
19     """
20     w_cls = 1 # for cls_scores
21     w_bbox = 1 # for offsets
22     total_loss = None
23     #####
24     # TODO: Implement the forward pass of Fast R-CNN. #
25     # A few key steps are outlined as follows: #
26     # i) Extract image feature. #
27     # ii) Perform RoI Align on proposals, then meanpool the feature in the #
28     #     spatial dimension. #
29     # iii) Pass the RoI feature through the shared-fc layer. Predict #
30     #     classification scores and box offsets. #
31     # iv) Assign the proposals with targets of each image. #
32     # v) Compute the cls_loss between the predicted class_prob and GT_class #
33     #     (For positive & negative proposals) #
34     #     Compute the bbox_loss between the offsets and GT_offsets #
35     #     (For positive proposals) #
36     #     Compute the total_loss which is formulated as: #
37     #     total_loss = w_cls*cls_loss + w_bbox*bbox_loss. #
38     #####
39     # Replace "pass" statement with your code
40     B, _, H, W = images.shape
41
42     # extract image feature
43     features = self.feats_extractor(images)
44

```

```

45     # perform RoI Pool & mean pool
46     boxes = torch.cat((proposal_batch_ids.unsqueeze(1), proposals), dim=-1)
47     roi_feat = torchvision.ops.roi_pool(features, boxes, (self.roi_output_w, self.roi_output_h))
48     roi_feat = roi_feat.mean(dim=[2, 3])
49
50     # forward heads, get predicted cls scores & offsets
51     shared_feat = self.shared_fc(roi_feat)
52     cls_scores = self.cls_head(shared_feat)
53     bbox_offsets = self.bbox_head(shared_feat)
54
55     # assign targets with proposals
56     pos_masks, neg_masks, GT_labels, GT_bboxes = [], [], [], []
57     for img_idx in range(B):
58         # get the positive/negative proposals and corresponding
59         # GT box & class label of this image
60         proposals_img = proposals[proposal_batch_ids == img_idx]
61         bboxes_img = bboxes[bbox_batch_ids == img_idx]
62         pos_mask, neg_mask, GT_label, GT_bbox = assign_label(proposals_img, bboxes_img, self.num_classes
63         )
64         pos_masks.append(pos_mask)
65         neg_masks.append(neg_mask)
66         GT_labels.append(GT_label)
67         GT_bboxes.append(GT_bbox)
68
69     # compute loss
70     cls_loss = 0
71     bbox_loss = 0
72     for img_idx in range(B):
73         pos_mask = pos_masks[img_idx]
74         neg_mask = neg_masks[img_idx]
75         GT_label = GT_labels[img_idx]
76         GT_bbox = GT_bboxes[img_idx]
77         proposals_img = proposals[proposal_batch_ids == img_idx]
78         cls_scores_img = cls_scores[proposal_batch_ids == img_idx]
79         bbox_offsets_img = bbox_offsets[proposal_batch_ids == img_idx]
80         cls_loss += ClsScoreRegression(cls_scores_img[pos_mask | neg_mask], GT_label[pos_mask | neg_mask
81         ], B)
82         bbox_loss += BboxRegression(bbox_offsets_img[pos_mask], compute_offsets(proposals_img[pos_mask],
83         GT_bbox), B)
84     total_loss = w_cls * cls_loss + w_bbox * bbox_loss
85     #####
86     #                                     #
87     #####
88     return total_loss

```

1.5 Task5

完成 `utils.py` 的 `generate_proposal` 函数。(完成代码即可，不用在报告中写文字说明)

代码如下：

```

1     def generate_proposal(anchors, offsets):
2         """

```

```

3 Proposal generator.
4
5 Inputs:
6 - anchors: Anchor boxes, of shape (M, 4). Anchors are represented
7 by the coordinates of their top-left and bottom-right corners.
8 - offsets: Transformations of shape (M, 4) that will be used to
9 convert anchor boxes into region proposals. The transformation
10 offsets[m] = (tx, ty, tw, th) will be applied to the anchor
11 anchors[m].
12
13 Outputs:
14 - proposals: Region proposals of shape (M, 4), represented by the
15 coordinates of their top-left and bottom-right corners. Applying the
16 transform offsets[m] to the anchor[m] should give the
17 proposal proposals[m].
18
19 """
20 proposals = None
21 #####
22 # TODO: Given anchor coordinates and the proposed offset for each anchor, #
23 # compute the proposal coordinates using the transformation formulas above. #
24 #####
25 # Replace "pass" statement with your code
26 xy_offsets = offsets[:, :2]
27 wh_offsets = offsets[:, 2:4]
28 proposals_minus = torch.exp(wh_offsets) * (anchors[:, 2:4] - anchors[:, :2])
29 proposals_plus = xy_offsets * (anchors[:, 2:4] - anchors[:, :2]) * 2 + (anchors[:, :2] + anchors[:,
30 2:4])
31 proposals = torch.cat(((proposals_plus - proposals_minus) / 2, (proposals_plus + proposals_minus) / 2)
32 , dim=-1)
33 #####
34 # END OF YOUR CODE #
35 #####
36
37 return proposals

```

1.6 Task6

完成 model.py 的 inference 函数。(完成代码即可，不用在报告中写文字说明)

代码如下：

```

1 def inference(self, images, proposals, proposal_batch_ids, thresh=0.5, nms_thresh=0.7):
2     """
3     Inference-time forward pass for our two-stage Faster R-CNN detector
4
5     Inputs:
6     - images: Tensor of shape (B, 3, H, W) giving input images
7     - proposals: Tensor of shape (M, 4) giving the proposals for input images,
8     where M is the total number of proposals in the batch
9     - proposal_batch_ids: Tensor of shape (M, ) giving the index of the image
10    that each proposals belongs to
11    - thresh: Threshold value on confidence probability. HINT: You can convert the

```

```

12     classification score to probability using a softmax nonlinearity.
13     - nms_thresh: IoU threshold for NMS
14
15     We can output a variable number of predicted boxes per input image.
16     In particular we assume that the input images[i] gives rise to P_i final
17     predicted boxes.
18
19     Outputs:
20     - final_proposals: List of length (B,) where final_proposals[i] is a Tensor
21     of shape (P_i, 4) giving the coordinates of the final predicted boxes for
22     the input images[i]
23     - final_conf_probs: List of length (B,) where final_conf_probs[i] is a
24     Tensor of shape (P_i, 1) giving the predicted probabilities that the boxes
25     in final_proposals[i] are objects (vs background)
26     - final_class: List of length (B,), where final_class[i] is an int64 Tensor
27     of shape (P_i, 1) giving the predicted category labels for each box in
28     final_proposals[i].
29     """
30     final_proposals, final_conf_probs, final_class = None, None, None
31     #####
32     # TODO: Predicting the final proposal coordinates `final_proposals`, #
33     # confidence scores `final_conf_probs`, and the class index `final_class`. #
34     # The overall steps are similar to the forward pass, but now you cannot #
35     # decide the activated nor negative proposals without GT boxes.      #
36     # You should apply post-processing (thresholding and NMS) to all proposals #
37     # and keep the final proposals.                                         #
38     #####
39     # Replace "pass" statement with your code
40     B = images.shape[0]
41
42     # extract image feature
43     features = self.feat_extractor(images)
44
45     # perform RoI Pool & mean pool
46     boxes = torch.cat((proposal_batch_ids.unsqueeze(1), proposals), dim=-1)
47     roi_feat = torchvision.ops.roi_pool(features, boxes, (self.roi_output_w, self.roi_output_h))
48     roi_feat = roi_feat.mean(dim=[2, 3])
49
50     # forward heads, get predicted cls scores & offsets
51     shared_feat = self.shared_fc(roi_feat)
52     cls_scores = self.cls_head(shared_feat)
53     bbox_offsets = self.bbox_head(shared_feat)
54
55     # get predicted boxes & class label & confidence probability
56     conf_probs = torch.softmax(cls_scores, dim=-1)
57     pred_boxes = generate_proposal(proposals, bbox_offsets)
58
59     final_proposals = []
60     final_conf_probs = []
61     final_class = []
62     # post-process to get final predictions
63     for img_idx in range(B):
64
65         # filter by threshold

```



```

66     img_proposals = pred_boxes[proposal_batch_ids == img_idx]
67     img_conf_probs = conf_probs[proposal_batch_ids == img_idx]
68     img_cls_scores = cls_scores[proposal_batch_ids == img_idx]
69     keep = img_conf_probs[:, :self.num_classes].max(dim=1).values > thresh
70     img_proposals = img_proposals[keep]
71     img_conf_probs = img_conf_probs[keep]
72     img_cls_scores = img_cls_scores[keep]
73     conf_values, pred_classes = img_conf_probs[:, :self.num_classes].max(dim=1)
74
75     # nms
76     keep_idx = torchvision.ops.nms(img_proposals, conf_values, nms_thresh)
77     final_proposals.append(img_proposals[keep_idx])
78     final_conf_probs.append(conf_values[keep_idx].unsqueeze(1))
79     final_class.append(pred_classes[keep_idx].unsqueeze(1))
80
81     #####
82     #                               END OF YOUR CODE                               #
83     #####
84     return final_proposals, final_conf_probs, final_class

```

1.7 Task7

完成过拟合实验，在报告中给出训练损失曲线和测试样本可视化。

训练损失函数曲线和测试样本可视化如下：





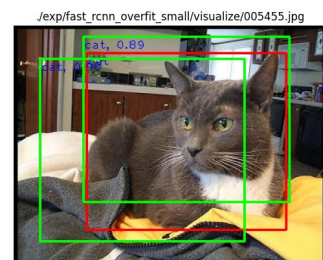






图 1: 训练损失曲线



注意到，训练损失函数曲线和“car”的测试样本与说明文档中的示例几乎完全一致，说明实验结果符合预期。

1.8 Task8

完成最终实验，在报告中给出训练损失曲线和评测情况。

训练损失曲线和评测情况如下：






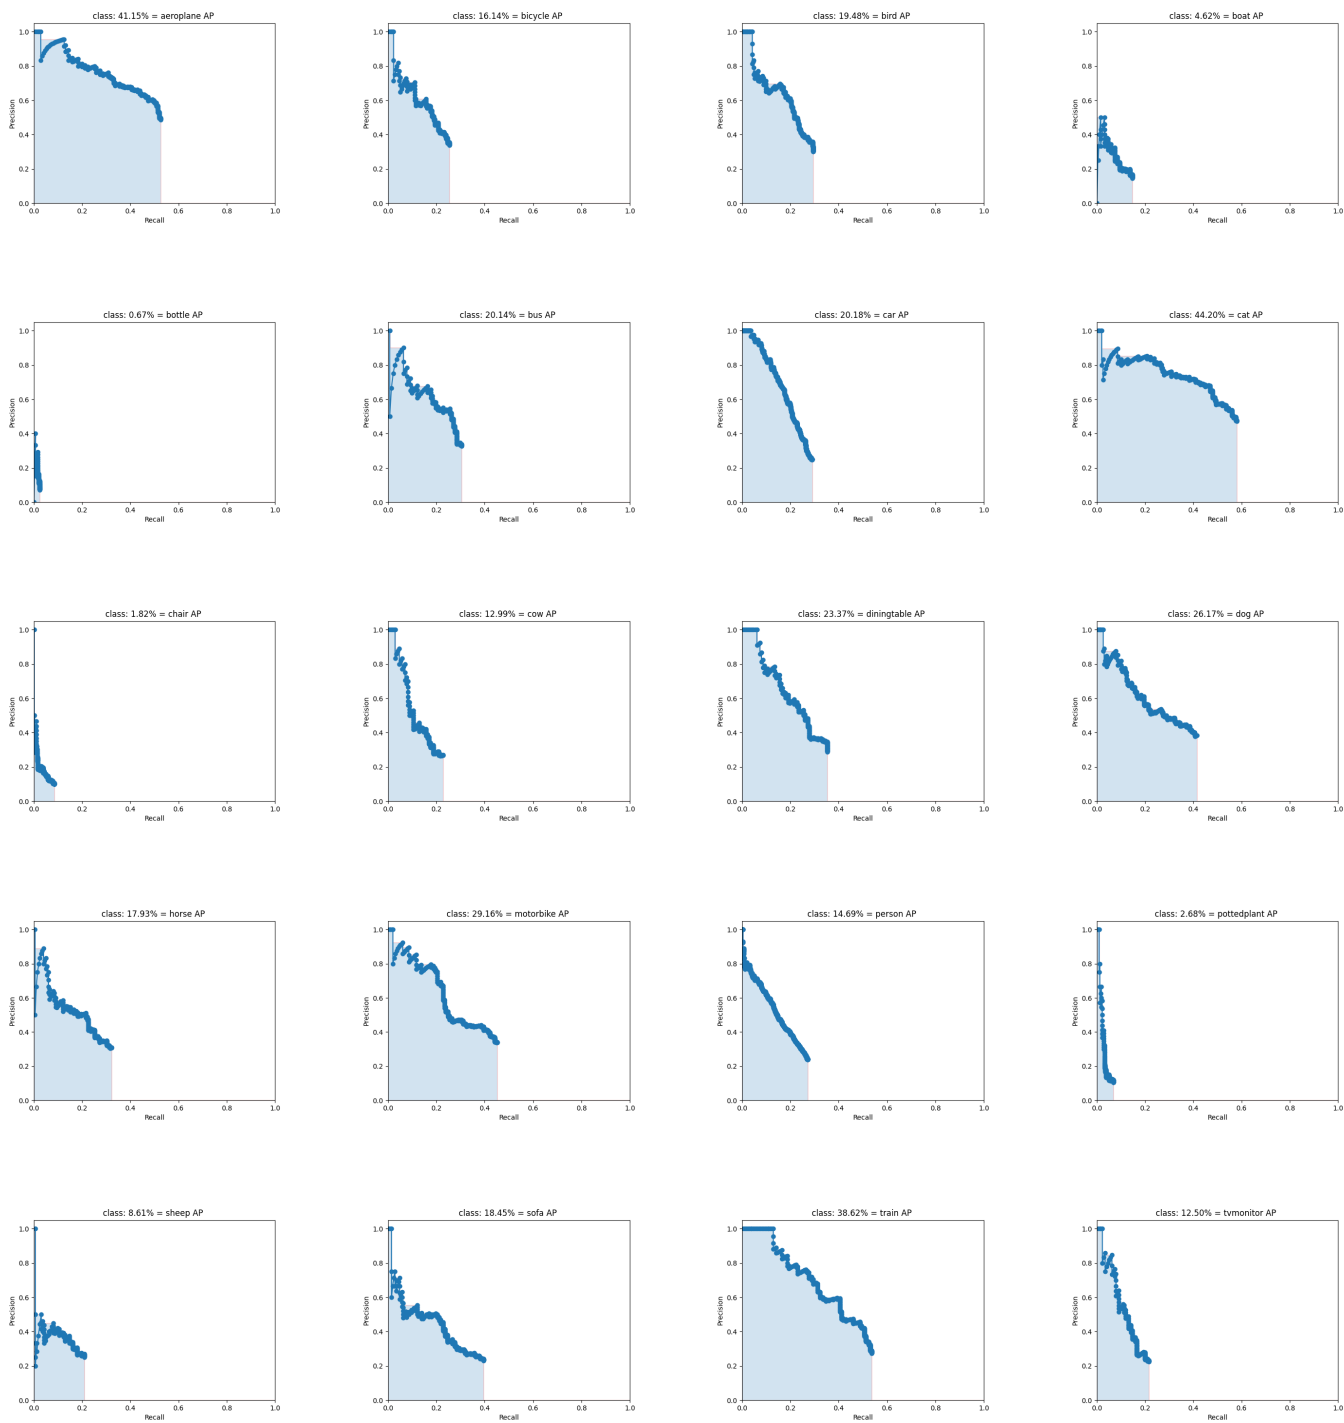


图 2: 最终训练损失曲线





受限于算力和时间，最终评测得到的 mAP 为 18.68%，与说明文档中的预期（18% 左右）相符。